Карпейкин В.А. ББМО-02-23 Номер 10

Клонирование репозитория

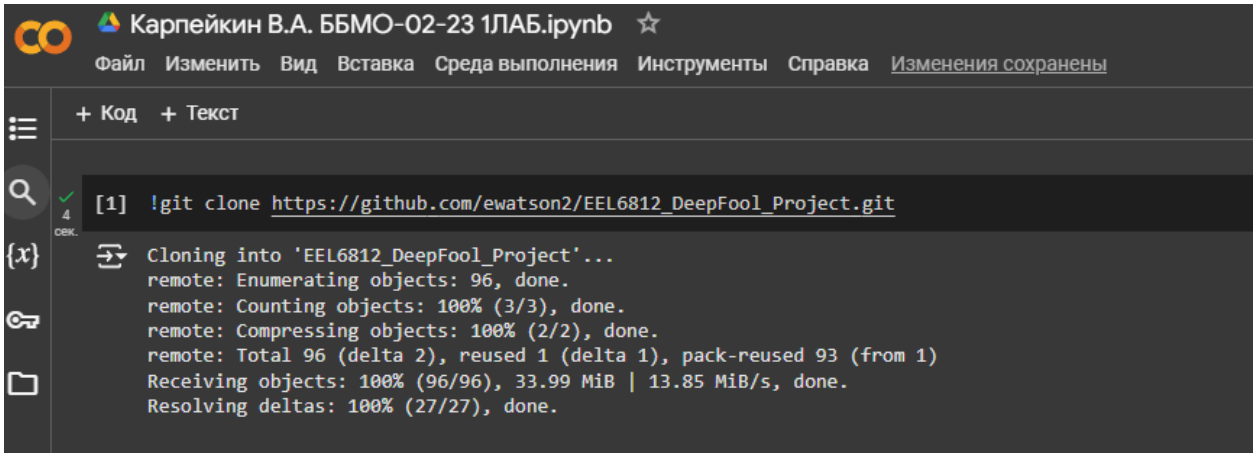

Смена директории и импорт библиотек

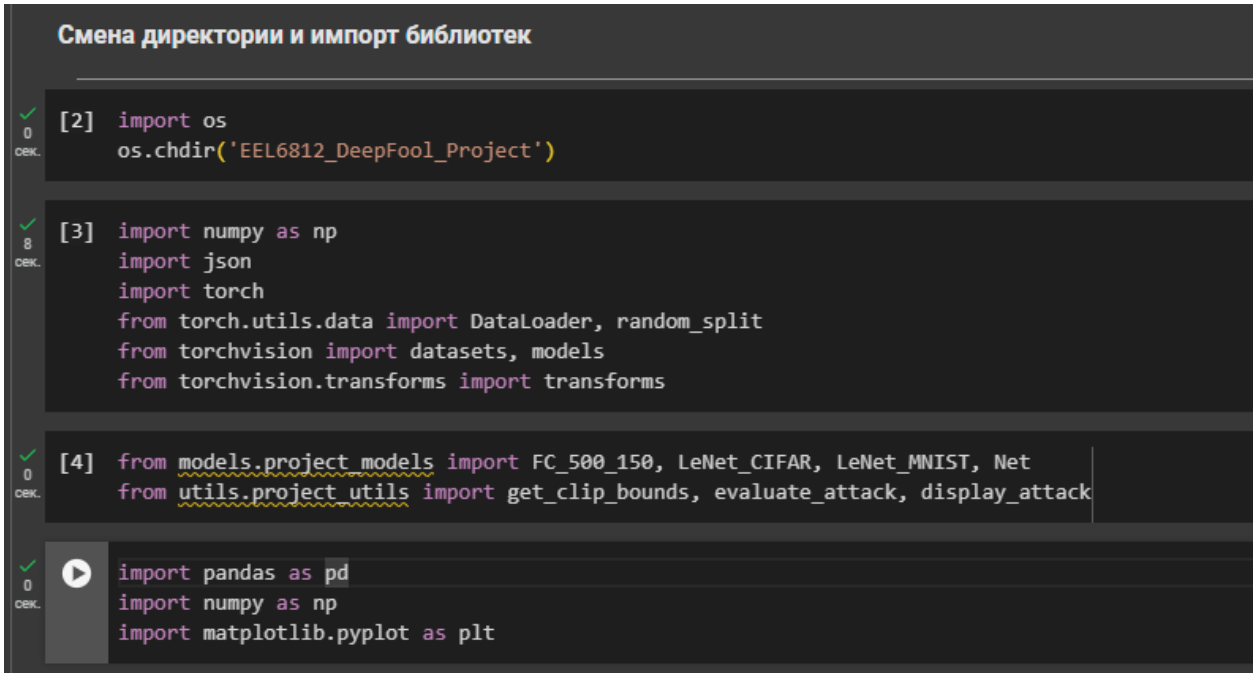

Установка случайного значения – номер в списке группы «7»

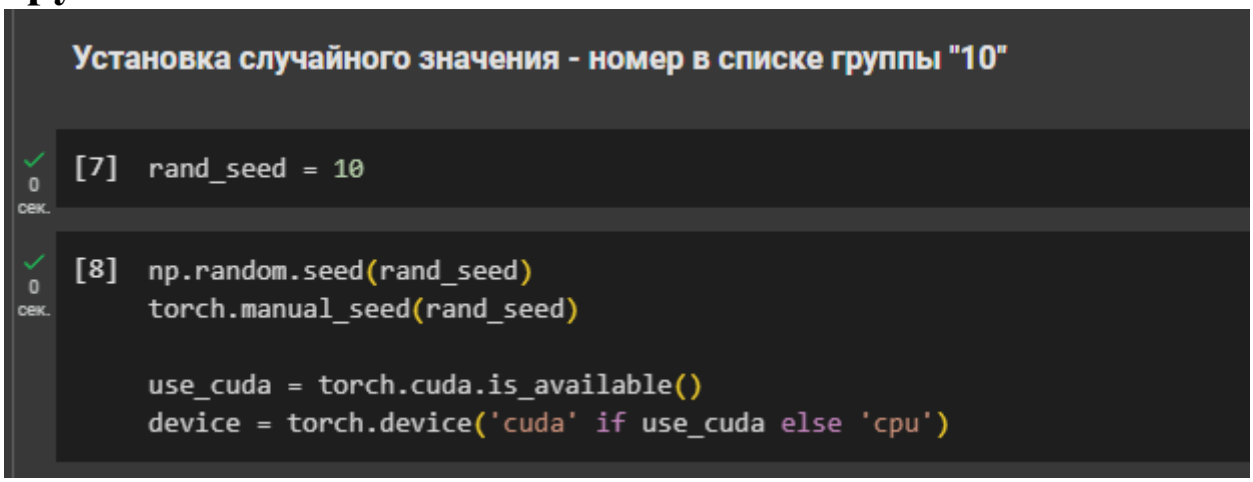

Загрузка датасета MNIST

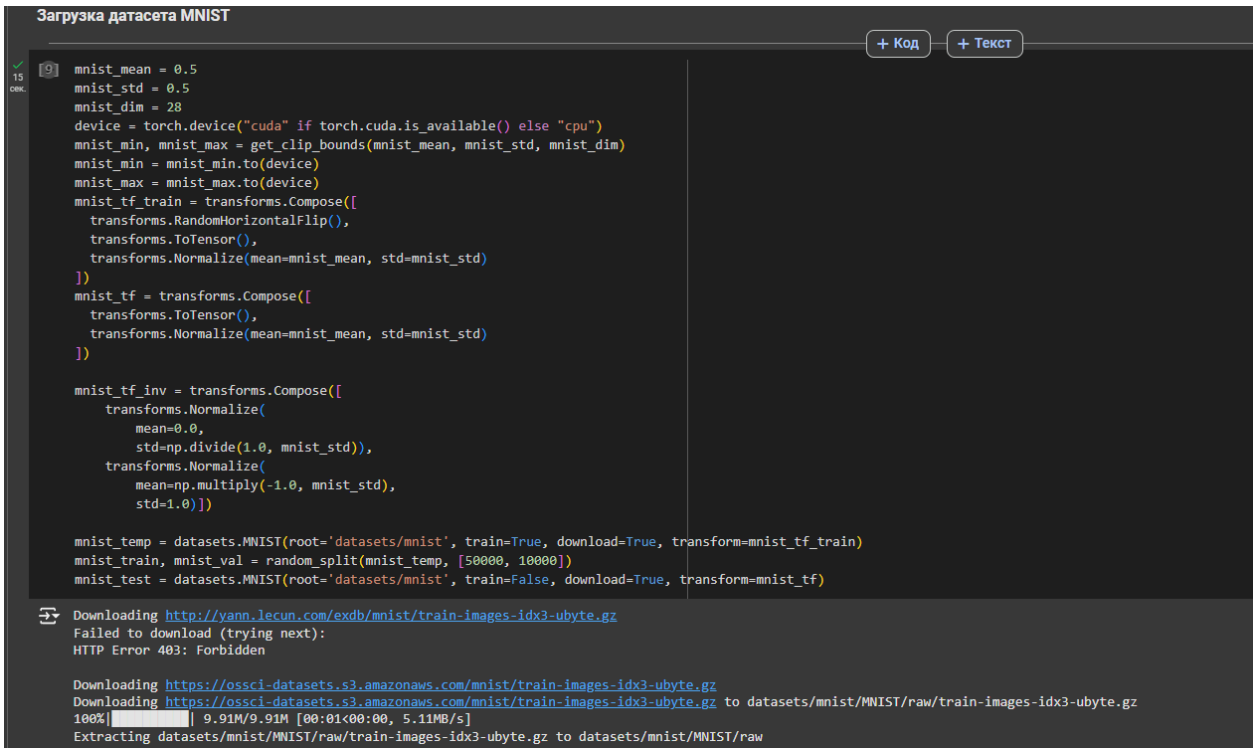

Загрузка датасета CIFAR-10

Загрузка датасета CIFAR-10

```
cifar_mean = [0.491, 0.482, 0.447]
cifar_std = [0.202, 0.199, 0.201]
cifar_dim = 32
cifar_min, cifar_max = get_clip_bounds(cifar_mean, cifar_std, cifar_dim)
cifar_min = cifar_min.to(device)
cifar_max = cifar_max.to(device)
cifar_tf_train = transforms.Compose([
  transforms.RandomCrop(size=cifar_dim, padding=4),
  transforms.RandomHorizontalFlip(),
  transforms.ToTensor(),
  transforms.Normalize(mean=cifar_mean, std=cifar_std)
])
cifar_tf = transforms.Compose([
  transforms.ToTensor(),
  transforms.Normalize(mean=cifar_mean, std=cifar_std)
])

cifar_tf_inv = transforms.Compose([
    transforms.Normalize(
        mean=[0.0, 0.0, 0.0],
        std=np.divide(1.0, cifar_std)),
    transforms.Normalize(
        mean=np.multiply(-1.0, cifar_mean),
        std=[1.0, 1.0, 1.0])])

cifar_temp = datasets.CIFAR10(root='datasets/cifar-10', train=True, download=True, transform=cifar_tf_train)
cifar_train, cifar_val = random_split(cifar_temp, [40000, 10000])
cifar_test = datasets.CIFAR10(root='datasets/cifar-10', train=False, download=True, transform=cifar_tf)
cifar_classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

```
Downloading https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz to datasets/cifar-10/cifar-10-python.tar.gz
100%|| 170M/170M [00:12<00:00, 13.1MB/s]
Extracting datasets/cifar-10/cifar-10-python.tar.gz to datasets/cifar-10
Files already downloaded and verified
```

Настройка DataLoader

Настройка DataLoader

```
[11] batch_size = 64
     workers = 4
     mnist_loader_train = DataLoader(mnist_train, batch_size=batch_size, shuffle=True, num_workers=workers)
     mnist_loader_val = DataLoader(mnist_val, batch_size=batch_size, shuffle=False, num_workers=workers)
     mnist_loader_test = DataLoader(mnist_test, batch_size=batch_size, shuffle=False, num_workers=workers)
     cifar_loader_train = DataLoader(cifar_train, batch_size=batch_size, shuffle=True, num_workers=workers)
     cifar_loader_val = DataLoader(cifar_val, batch_size=batch_size, shuffle=False, num_workers=workers)
     cifar_loader_test = DataLoader(cifar_test, batch_size=batch_size, shuffle=False, num_workers=workers)
```

```
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will create 4 worker processes in total. Our suggested max number of worker in current system is 2, which is smaller than what th
  warnings.warn(
```

FGSM атака

Стойкость к атаке моделей LeNet, FC на датасете MNIST и стойкость к атаке моделей Network-In-Network, LeNet на датасете CIFAR-10

LeNet MNIST

fgsm_eps = 0.001

```
<ipython-input-27-a421ced483ef>:2: FutureWarning: You are using
  model.load_state_dict(torch.load('weights/clean/mnist_lenet.p
/usr/local/lib/python3.10/dist-packages/torch/utils/data/datalo
  warnings.warn(
Точность до атаки: 98.34%
/usr/local/lib/python3.10/dist-packages/torch/utils/data/datalo
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 1.69%
FGSM Robustness : 8.06e-04
FGSM Time (All Images) : 1.05 s
FGSM Time (Per Image) : 104.57 us
```

fgsm_eps = 0.02

```
<ipython-input-29-a421ced483ef>:2: FutureWarning: You are using
  model.load_state_dict(torch.load('weights/clean/mnist_lenet.pt
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloa
  warnings.warn(
Точность до атаки: 98.34%
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloa
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 2.56%
FGSM Robustness : 1.59e-02
FGSM Time (All Images) : 0.95 s
FGSM Time (Per Image) : 95.34 us
```

fgsm_eps = 0.5

```
<ipython-input-31-a421ced483ef>:2: FutureWarnin
  model.load_state_dict(torch.load('weights/cle
/usr/local/lib/python3.10/dist-packages/torch/u
  warnings.warn(
Точность до атаки: 98.34%
/usr/local/lib/python3.10/dist-packages/torch/u
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 82.92%
FGSM Robustness : 3.83e-01
FGSM Time (All Images) : 1.00 s
FGSM Time (Per Image) : 99.73 us
```

fgsm_eps = 0.9

```
<ipython-input-33-a421ced483ef>:2: FutureWarning: You
  model.load_state_dict(torch.load('weights/clean/mnis
/usr/local/lib/python3.10/dist-packages/torch/utils/da
  warnings.warn(
Точность до атаки: 98.34%
/usr/local/lib/python3.10/dist-packages/torch/utils/da
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 93.80%
FGSM Robustness : 6.81e-01
FGSM Time (All Images) : 1.04 s
FGSM Time (Per Image) : 103.73 us
```

fgsm_eps = 10

```
<ipython-input-35-a421ced483ef>:2: Future
  model.load_state_dict(torch.load('weigh
/usr/local/lib/python3.10/dist-packages/t
  warnings.warn(
Точность до атаки: 98.34%
FGSM Test Error : 94.15%
FGSM Robustness : 1.46e+00
FGSM Time (All Images) : 1.41 s
FGSM Time (Per Image) : 141.28 us
```

График

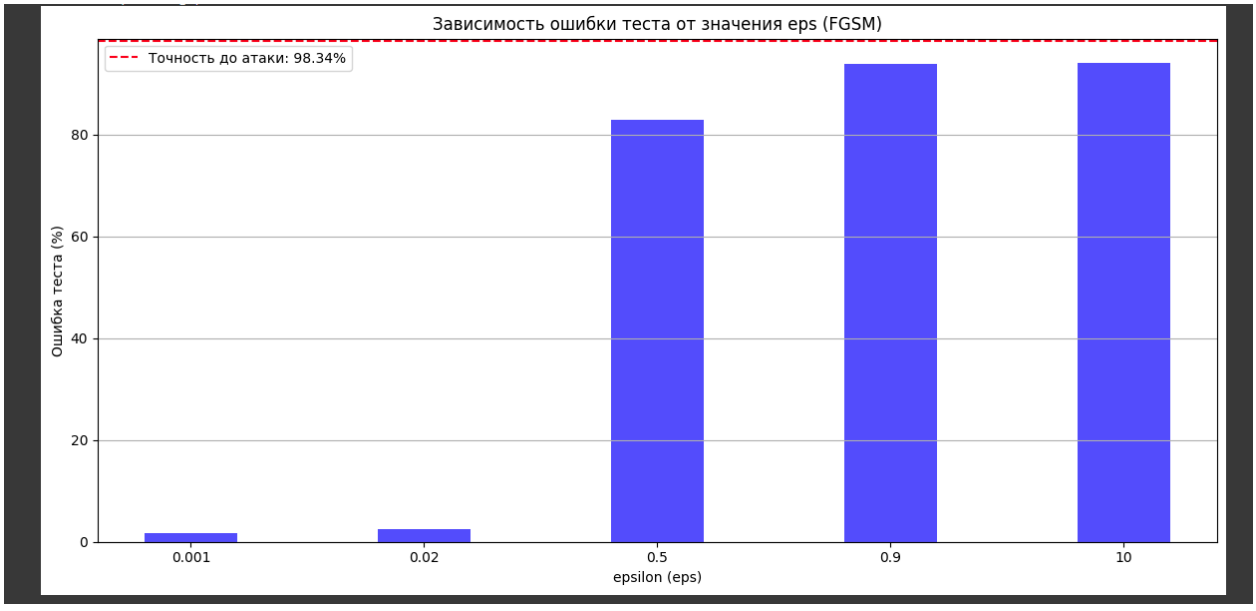

FC MNIST

fgsm_eps = 0.001

```
<ipython-input-37-86c4b3caf57a>:2: FutureWa
  model.load_state_dict(torch.load('weights
/usr/local/lib/python3.10/dist-packages/tor
  warnings.warn(
Точность до атаки: 97.03%
/usr/local/lib/python3.10/dist-packages/tor
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 3.07%
FGSM Robustness : 8.08e-04
FGSM Time (All Images) : 0.67 s
FGSM Time (Per Image) : 67.24 us
```

fgsm_eps = 0.02

```
<ipython-input-39-86c4b3caf57a>:2: FutureWarning: You are using `torch.load`
  model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: U
  warnings.warn(
Точность до атаки: 97.03%
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: U
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 5.54%
FGSM Robustness : 1.60e-02
FGSM Time (All Images) : 0.64 s
FGSM Time (Per Image) : 63.89 us
```

fgsm_eps = 0.5

```
<ipython-input-41-86c4b3caf57a>:2: Future
  model.load_state_dict(torch.load('weigh
/usr/local/lib/python3.10/dist-packages/t
  warnings.warn(
Точность до атаки: 97.03%
/usr/local/lib/python3.10/dist-packages/t
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 99.21%
FGSM Robustness : 3.86e-01
FGSM Time (All Images) : 0.63 s
FGSM Time (Per Image) : 63.05 us
```

fgsm_eps = 0.9

```
<ipython-input-43-86c4b3caf57a>:2: Futu
  model.load_state_dict(torch.load('wei
/usr/local/lib/python3.10/dist-packages
  warnings.warn(
Точность до атаки: 97.03%
/usr/local/lib/python3.10/dist-packages
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 99.87%
FGSM Robustness : 6.86e-01
FGSM Time (All Images) : 0.63 s
FGSM Time (Per Image) : 62.78 us
```

fgsm_eps = 10

```
<ipython-input-45-86c4b3caf57a>:2: Fut
  model.load_state_dict(torch.load('we
/usr/local/lib/python3.10/dist-package
  warnings.warn(
Точность до атаки: 97.03%
FGSM Test Error : 99.87%
FGSM Robustness : 1.47e+00
FGSM Time (All Images) : 0.85 s
FGSM Time (Per Image) : 84.99 us
```

График

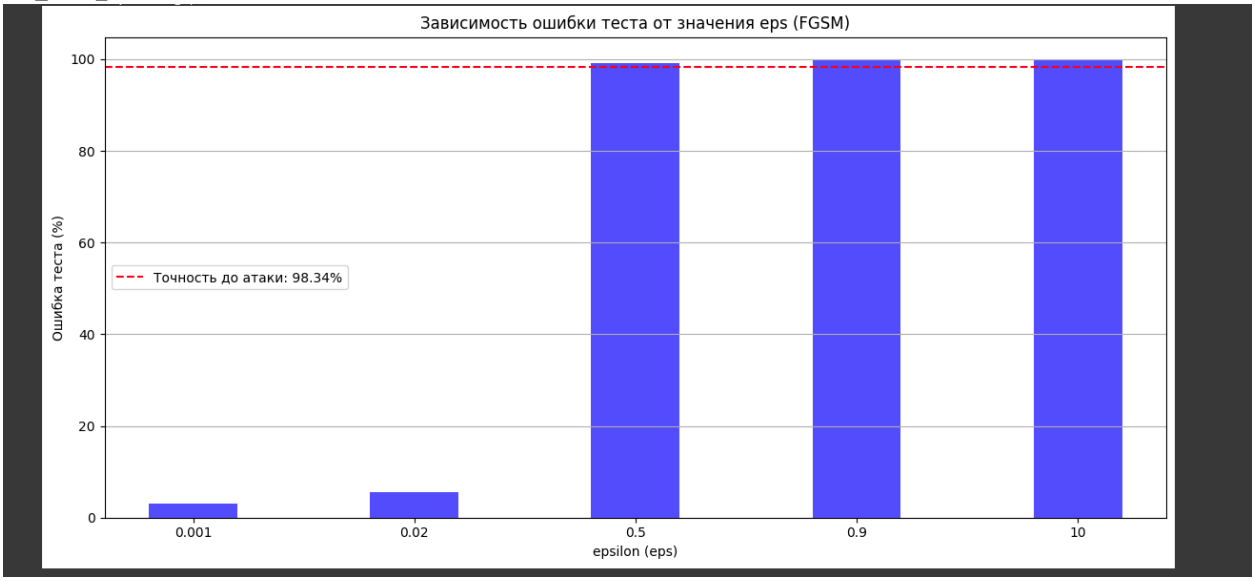

Network-In-Network CIFAR-10

fgsm_eps = 0.001

```
<ipython-input-47-412f7b94b9eb>:2: FutureWarning: You are u
  model.load_state_dict(torch.load('weights/clean/cifar_nir
/usr/local/lib/python3.10/dist-packages/torch/utils/data/da
  warnings.warn(
Точность до атаки: 90.72%
/usr/local/lib/python3.10/dist-packages/torch/utils/data/da
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 10.12%
FGSM Robustness : 8.92e-04
FGSM Time (All Images) : 1.29 s
FGSM Time (Per Image) : 129.21 us
```

fgsm_eps = 0.02

```
<ipython-input-50-412f7b94b9eb>:2: FutureW
  model.load_state_dict(torch.load('weight
/usr/local/lib/python3.10/dist-packages/to
  warnings.warn(
Точность до атаки: 90.72%
/usr/local/lib/python3.10/dist-packages/to
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 30.76%
FGSM Robustness : 1.78e-02
FGSM Time (All Images) : 1.68 s
FGSM Time (Per Image) : 167.74 us
```

fgsm_eps = 0.5

```
<ipython-input-53-412f7b94b9eb>:2: FutureWarning: You ar
  model.load_state_dict(torch.load('weights/clean/cifar_
/usr/local/lib/python3.10/dist-packages/torch/utils/data
  warnings.warn(
Точность до атаки: 90.72%
/usr/local/lib/python3.10/dist-packages/torch/utils/data
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 82.67%
FGSM Robustness : 4.40e-01
FGSM Time (All Images) : 1.54 s
FGSM Time (Per Image) : 153.98 us
```

fgsm_eps = 0.9

```
<ipython-input-56-412f7b94b9eb>:2: Future
  model.load_state_dict(torch.load('weigh
/usr/local/lib/python3.10/dist-packages/t
  warnings.warn(
Точность до атаки: 90.72%
/usr/local/lib/python3.10/dist-packages/t
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 84.62%
FGSM Robustness : 7.79e-01
FGSM Time (All Images) : 1.57 s
FGSM Time (Per Image) : 157.02 us
```

fgsm_eps = 10

```
<ipython-input-59-412f7b94b9eb>:2: Futur
  model.load_state_dict(torch.load('weig
/usr/local/lib/python3.10/dist-packages/
  warnings.warn(
Точность до атаки: 90.72%
FGSM Test Error : 87.50%
FGSM Robustness : 2.46e+00
FGSM Time (All Images) : 1.70 s
FGSM Time (Per Image) : 169.52 us
```

График

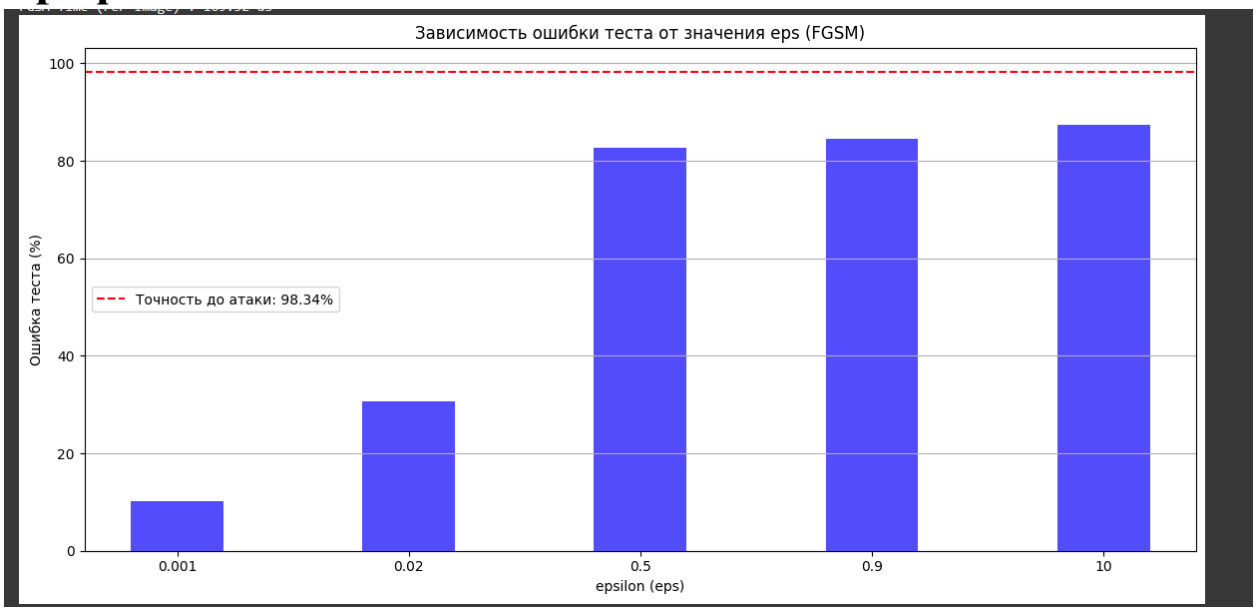

LeNet CIFAR-10

fgsm_eps = 0.001

```
<ipython-input-48-fe2d187d7de7>:2: FutureWarning
  model.load_state_dict(torch.load('weights/clea
/usr/local/lib/python3.10/dist-packages/torch/ut
  warnings.warn(
Точность до атаки: 78.66%
/usr/local/lib/python3.10/dist-packages/torch/ut
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 22.72%
FGSM Robustness : 8.92e-04
FGSM Time (All Images) : 1.41 s
FGSM Time (Per Image) : 140.76 us
```

fgsm_eps = 0.02

```
<ipython-input-51-fe2d187d7de7>:2: FutureWar
  model.load_state_dict(torch.load('weights/
/usr/local/lib/python3.10/dist-packages/torc
  warnings.warn(
Точность до атаки: 78.66%
/usr/local/lib/python3.10/dist-packages/torc
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 47.76%
FGSM Robustness : 1.78e-02
FGSM Time (All Images) : 1.29 s
FGSM Time (Per Image) : 128.99 us
```

fgsm_eps = 0.5

```
<ipython-input-54-fe2d187d7de7>:2: FutureWa
  model.load_state_dict(torch.load('weights
/usr/local/lib/python3.10/dist-packages/tor
  warnings.warn(
Точность до атаки: 78.66%
/usr/local/lib/python3.10/dist-packages/tor
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 95.17%
FGSM Robustness : 4.40e-01
FGSM Time (All Images) : 1.26 s
FGSM Time (Per Image) : 125.84 us
```

fgsm_eps = 0.9

```
<ipython-input-57-fe2d187d7de7>:2: FutureWarning:
  model.load_state_dict(torch.load('weights/clean/
/usr/local/lib/python3.10/dist-packages/torch/util
  warnings.warn(
Точность до атаки: 78.66%
/usr/local/lib/python3.10/dist-packages/torch/util
  warnings.warn(
FGSM Batches Complete : (157 / 157)
FGSM Test Error : 92.04%
FGSM Robustness : 7.80e-01
FGSM Time (All Images) : 1.33 s
FGSM Time (Per Image) : 133.41 us
```

fgsm_eps = 10

```
<ipython-input-60-fe2d187d7de7>:2: Futu
  model.load_state_dict(torch.load('wei
/usr/local/lib/python3.10/dist-packages
  warnings.warn(
Точность до атаки: 78.66%
FGSM Test Error : 89.90%
FGSM Robustness : 2.47e+00
FGSM Time (All Images) : 1.25 s
FGSM Time (Per Image) : 124.57 us
```

График

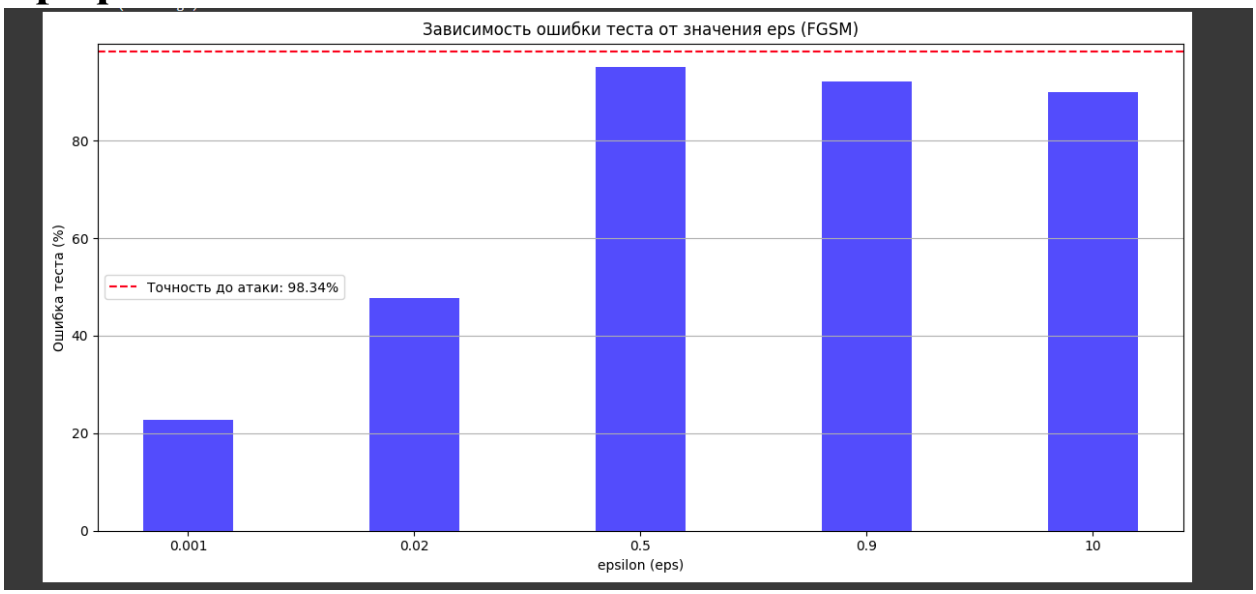

DeepFool атака

Стойкость к атаке моделей LeNet, FC на датасете MNIST и стойкость к атакае моделей Network-In-Network, LeNet на датасете CIFAR-10

LeNet MNIST

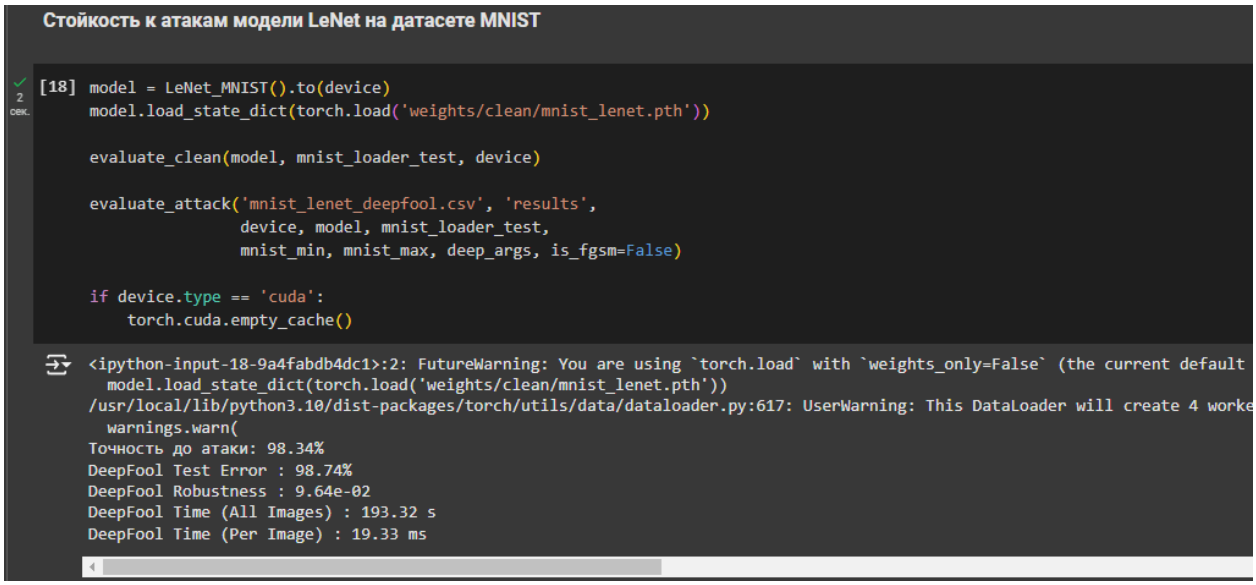

Стойкость к атакам модели LeNet на датасете MNIST

```
[18] model = LeNet_MNIST().to(device)
     model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))

     evaluate_clean(model, mnist_loader_test, device)

     evaluate_attack('mnist_lenet_deepfool.csv', 'results',
                     device, model, mnist_loader_test,
                     mnist_min, mnist_max, deep_args, is_fgsm=False)

     if device.type == 'cuda':
         torch.cuda.empty_cache()
```

```
<ipython-input-18-9a4fabdb4dc1>:2: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default
  model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: This DataLoader will create 4 worke
  warnings.warn(
Точность до атаки: 98.34%
DeepFool Test Error : 98.74%
DeepFool Robustness : 9.64e-02
DeepFool Time (All Images) : 193.32 s
DeepFool Time (Per Image) : 19.33 ms
```

FC MNIST

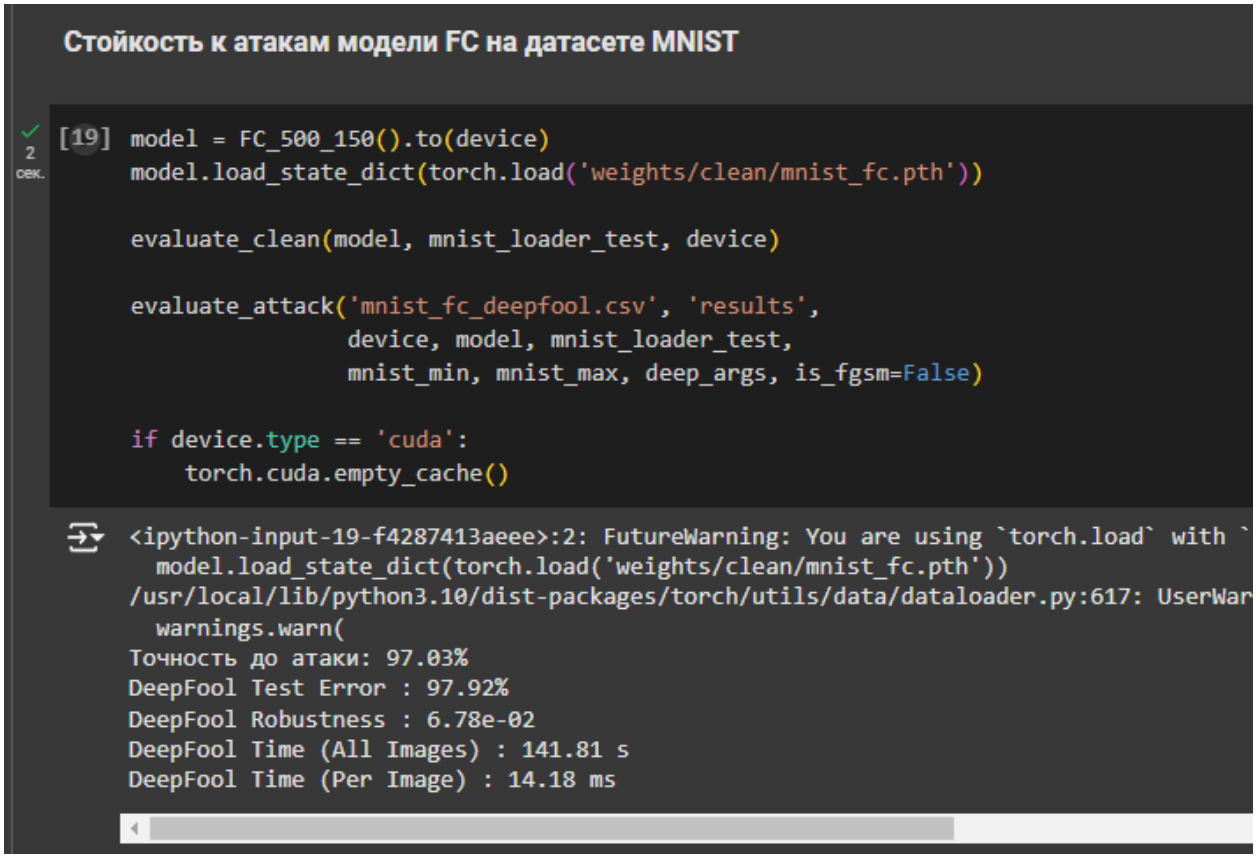

Стойкость к атакам модели FC на датасете MNIST

```
[19] model = FC_500_150().to(device)
     model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))

     evaluate_clean(model, mnist_loader_test, device)

     evaluate_attack('mnist_fc_deepfool.csv', 'results',
                     device, model, mnist_loader_test,
                     mnist_min, mnist_max, deep_args, is_fgsm=False)

     if device.type == 'cuda':
         torch.cuda.empty_cache()
```

```
<ipython-input-19-f4287413aeee>:2: FutureWarning: You are using `torch.load` with `
  model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWar
  warnings.warn(
Точность до атаки: 97.03%
DeepFool Test Error : 97.92%
DeepFool Robustness : 6.78e-02
DeepFool Time (All Images) : 141.81 s
DeepFool Time (Per Image) : 14.18 ms
```

Network-In-Network CIFAR-10

Стойкость к атакам модели Network-In-Network на датасете CIFAR-10

```
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))

evaluate_clean(model, cifar_loader_test, device)

evaluate_attack('cifar_nin_deepfool.csv', 'results',
                device, model, cifar_loader_test,
                cifar_min, cifar_max, deep_args, is_fgsm=False)

if device.type == 'cuda':
    torch.cuda.empty_cache()
```

```
<ipython-input-20-d39c82e071ac>:2: FutureWarning: You are using `torch.load` with `wei
  model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarnir
  warnings.warn(
Точность до атаки: 90.72%
DeepFool Test Error : 93.76%
DeepFool Robustness : 2.12e-02
DeepFool Time (All Images) : 185.12 s
DeepFool Time (Per Image) : 18.51 ms
```

LeNet CIFAR-10

Стойкость к атакам модели LeNet на датасете CIFAR-10

```
[21] model = LeNet_CIFAR().to(device)
     model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))

     evaluate_clean(model, cifar_loader_test, device)

     evaluate_attack('cifar_lenet_deepfool.csv', 'results',
                     device, model, cifar_loader_test,
                     cifar_min, cifar_max, deep_args, is_fgsm=False)

     if device.type == 'cuda':
         torch.cuda.empty_cache()
```

```
<ipython-input-21-71a3964ca979>:2: FutureWarning: You are using `torch.load` with `weights_
  model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
/usr/local/lib/python3.10/dist-packages/torch/utils/data/dataloader.py:617: UserWarning: Th
  warnings.warn(
Точность до атаки: 78.66%
DeepFool Test Error : 87.81%
DeepFool Robustness : 1.78e-02
DeepFool Time (All Images) : 73.27 s
DeepFool Time (Per Image) : 7.33 ms
```

Визуальное представление

LeNet MNIST

fgsm_eps = 0.001

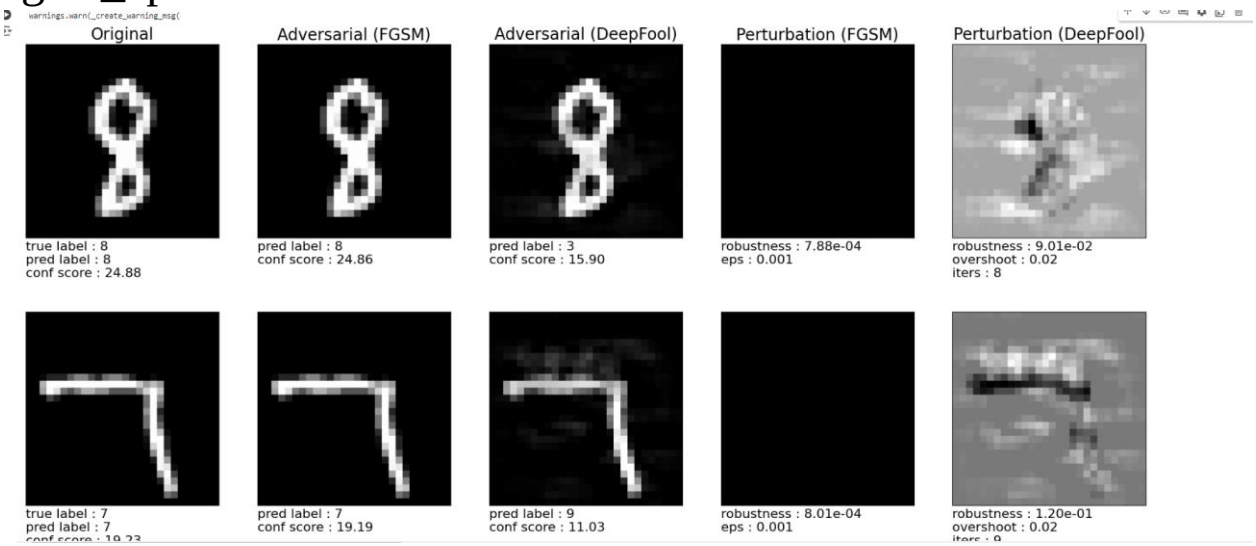

fgsm_eps = 0.02

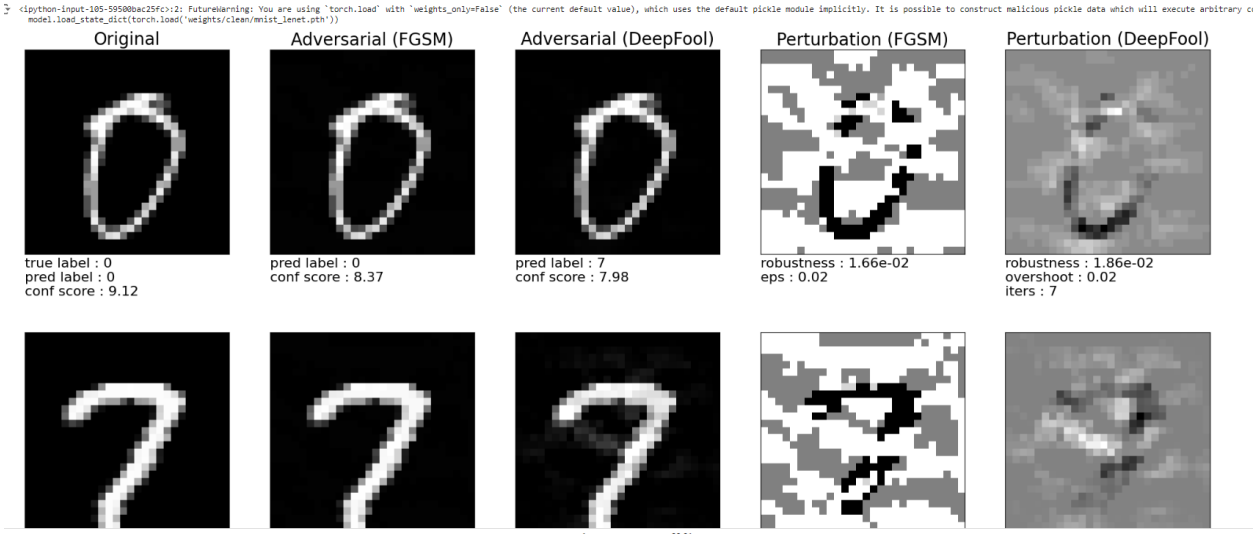

fgsm_eps = 0.5

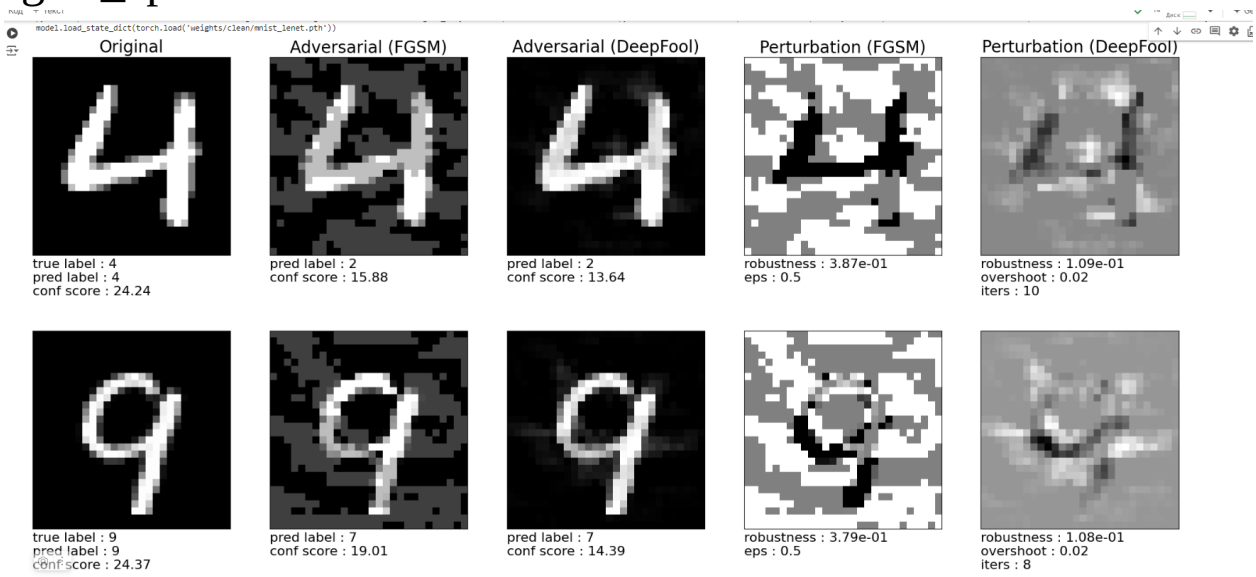

fgsm_eps = 0.9

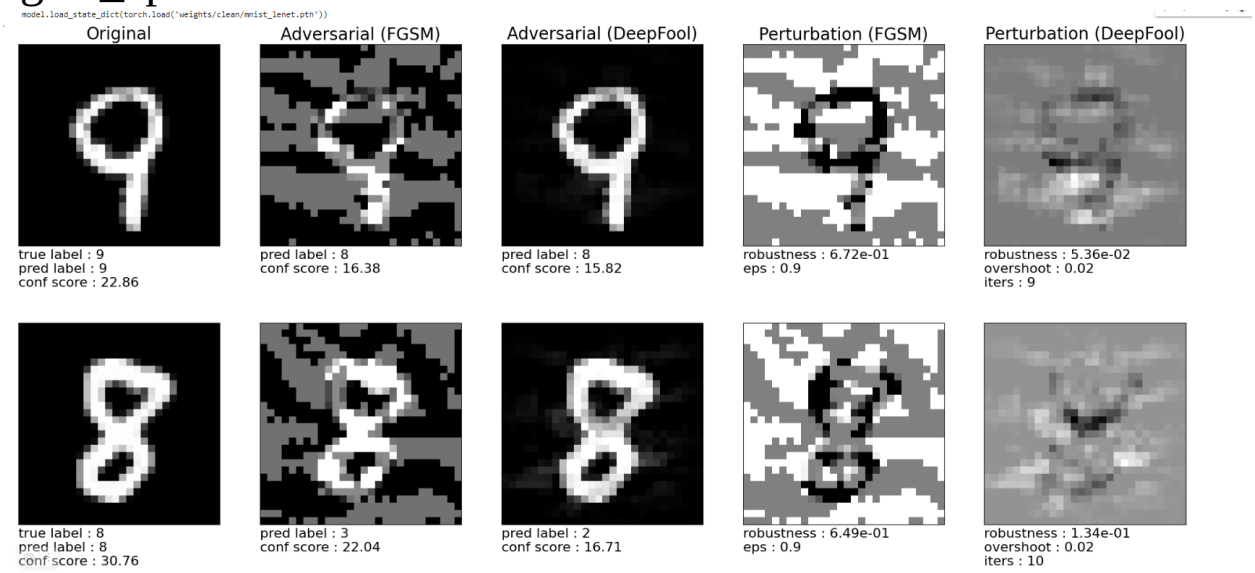

fgsm_eps = 10

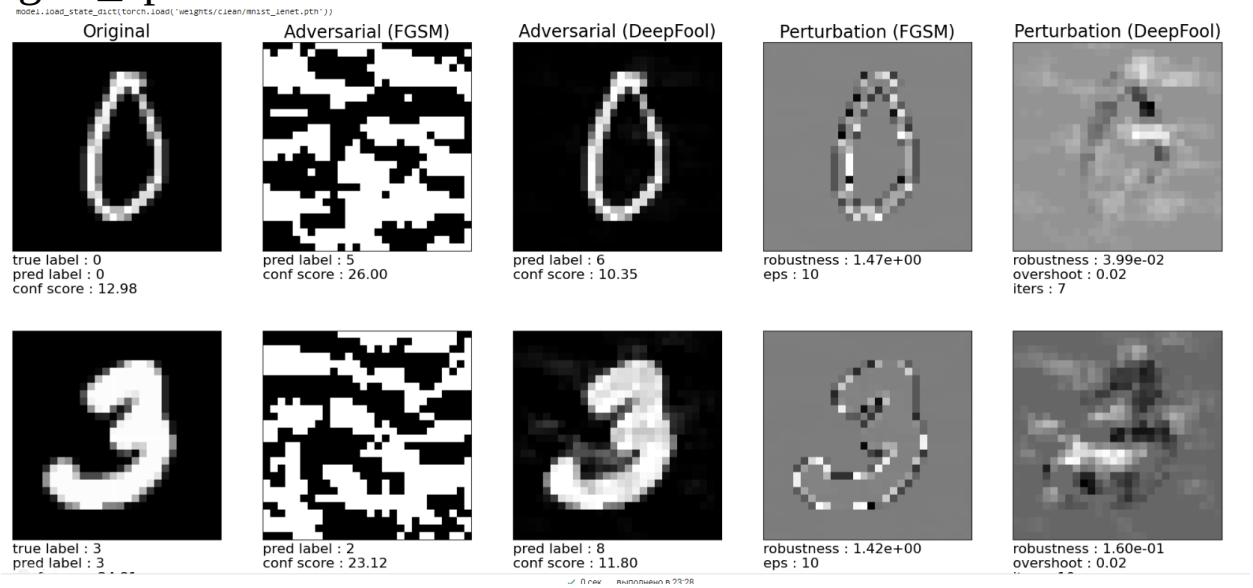

FC MNIST

fgsm_eps = 0.001

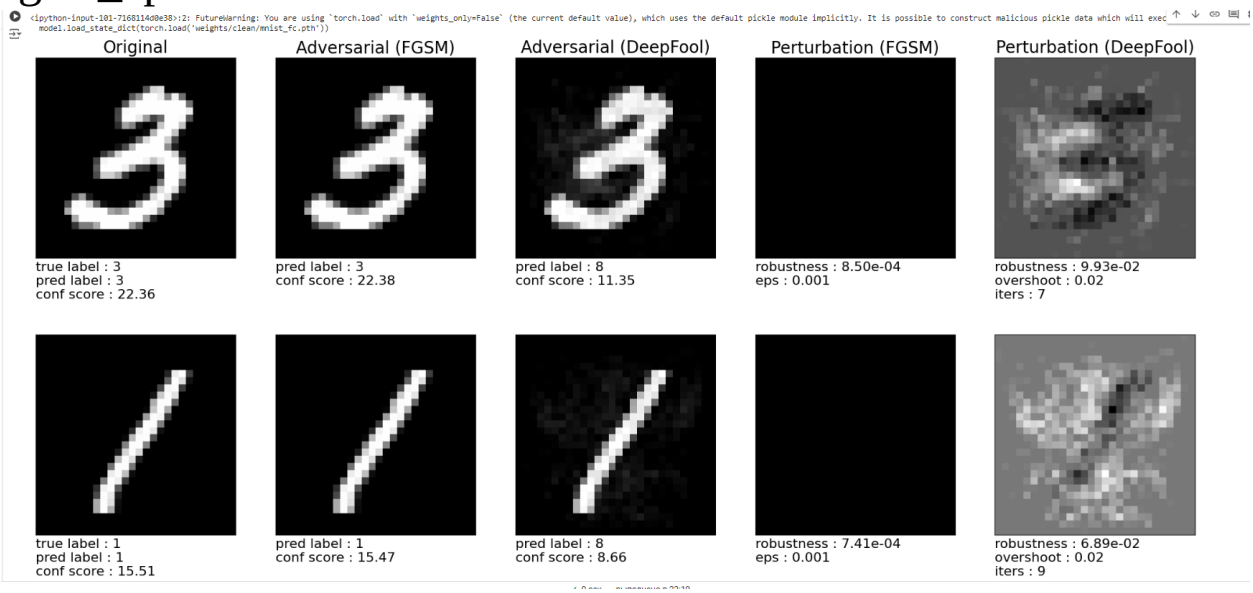

fgsm_eps = 0.02

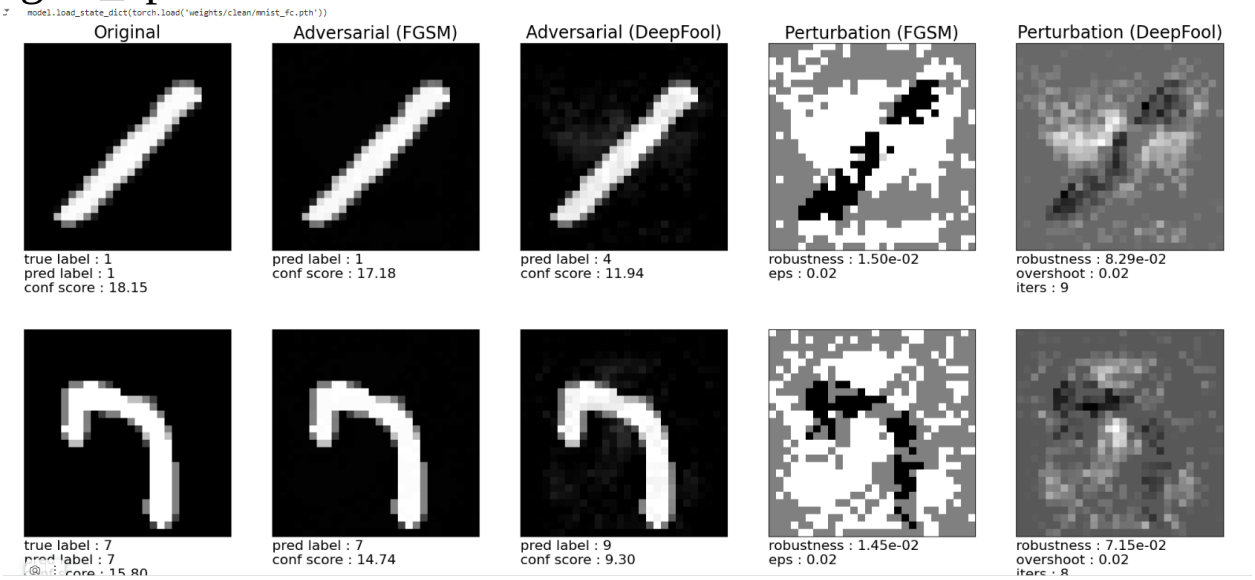

fgsm_eps = 0.5

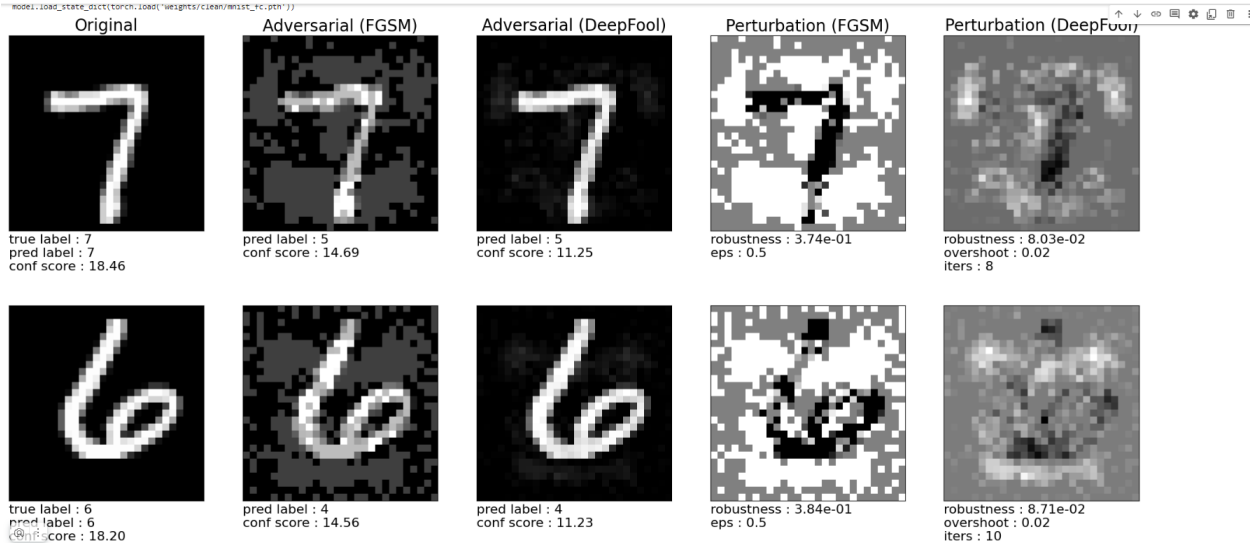

fgsm_eps = 0.9

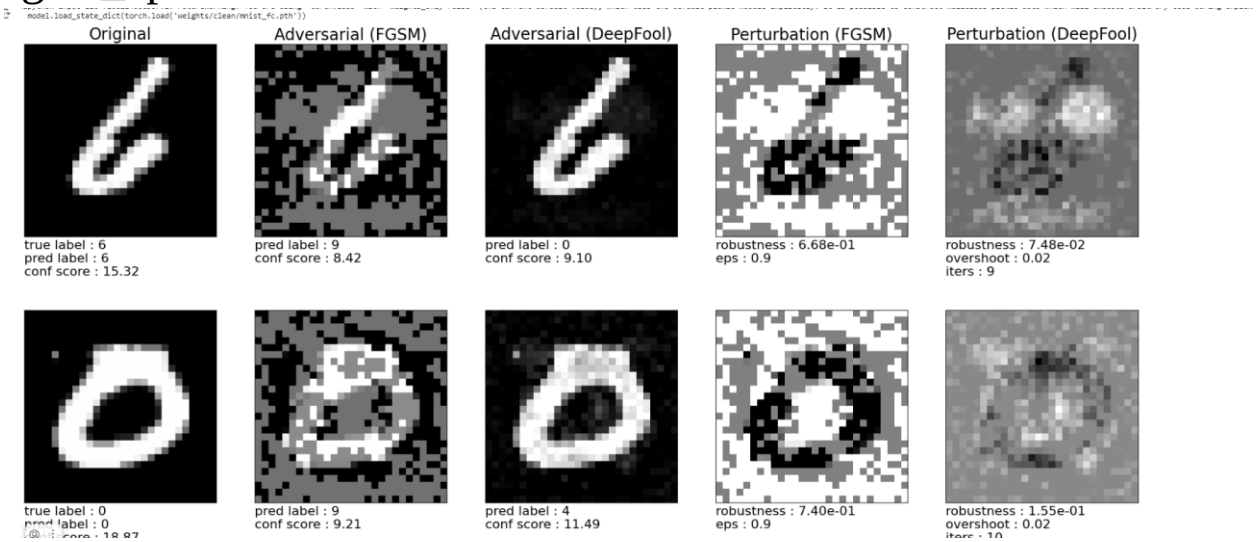

fgsm_eps = 10

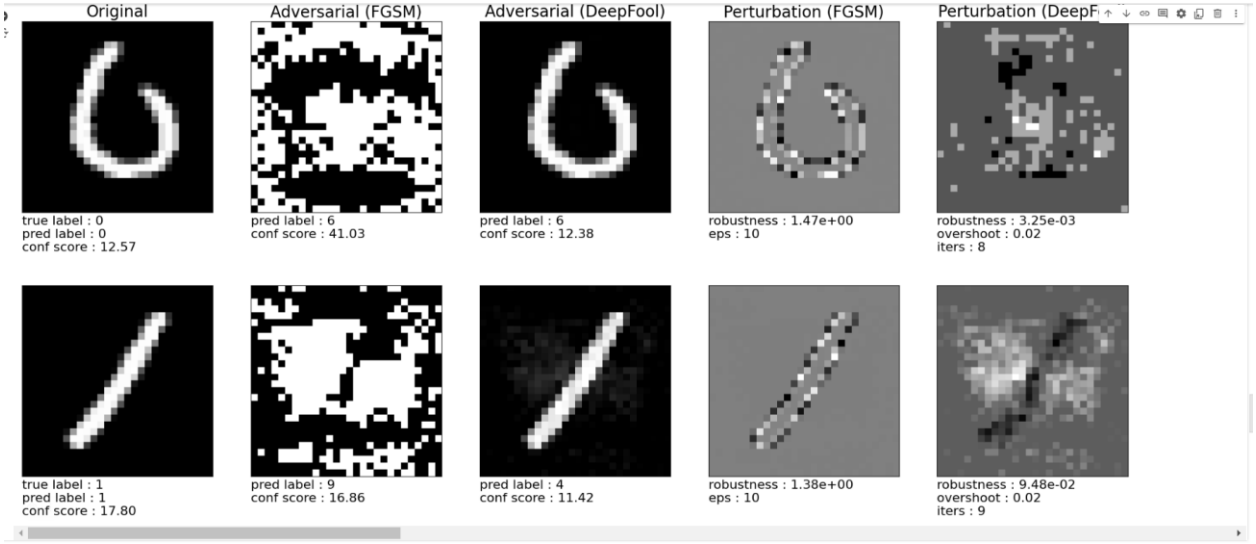

Network-In-Network CIFAR-10

fgsm_eps = 0.001

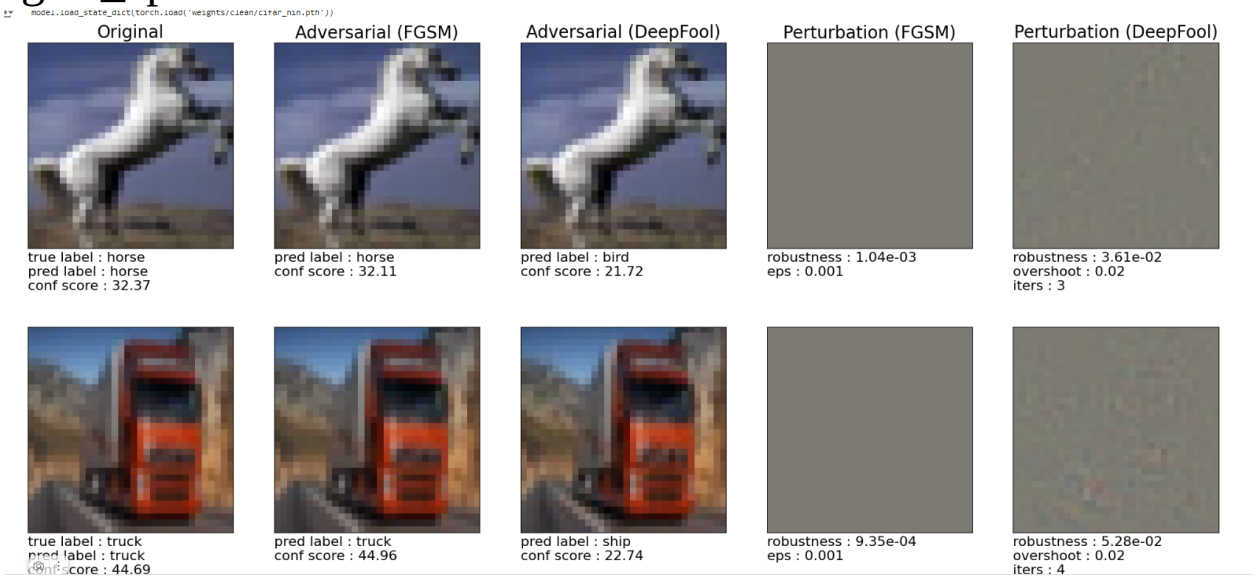

fgsm_eps = 0.02

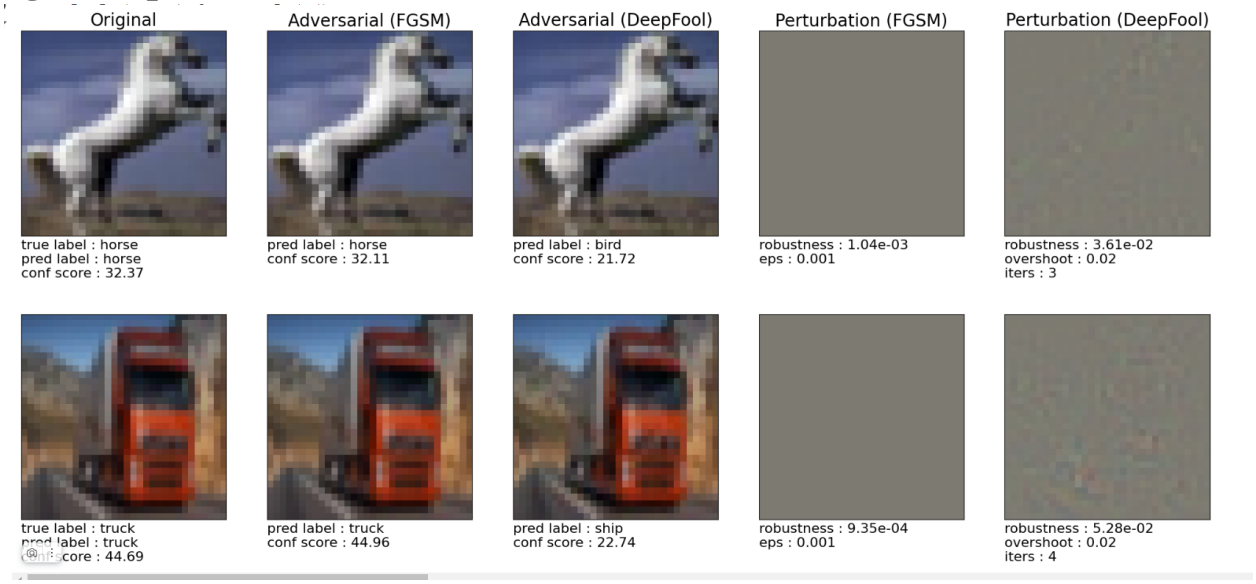

fgsm_eps = 0.5

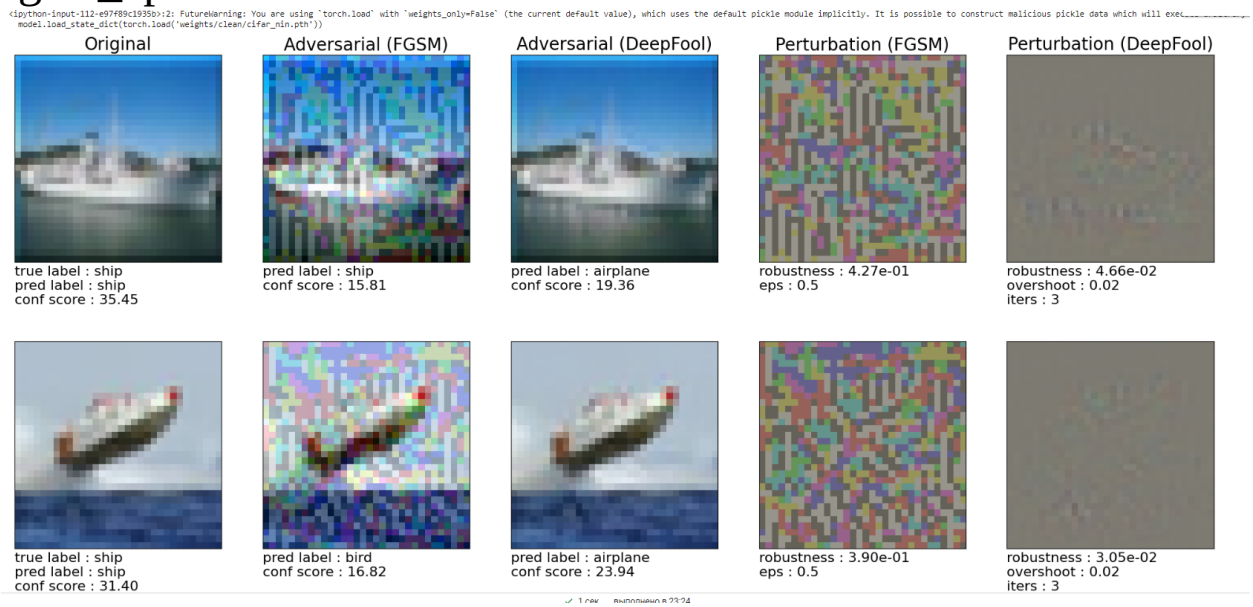

fgsm_eps = 0.9

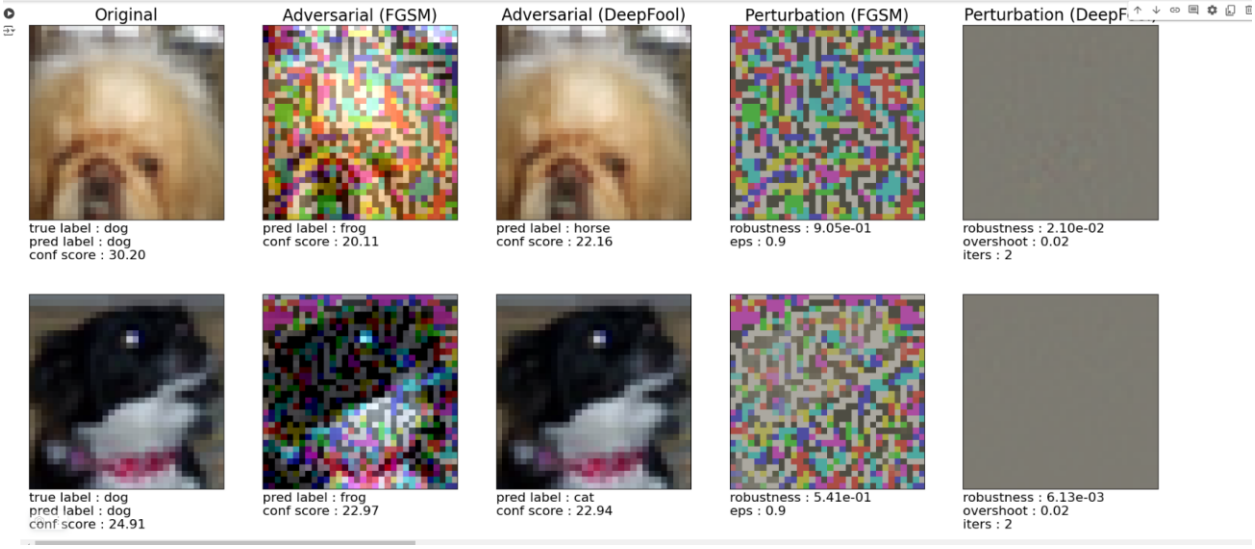

fgsm_eps = 10

LeNet CIFAR-10

fgsm_eps = 0.001

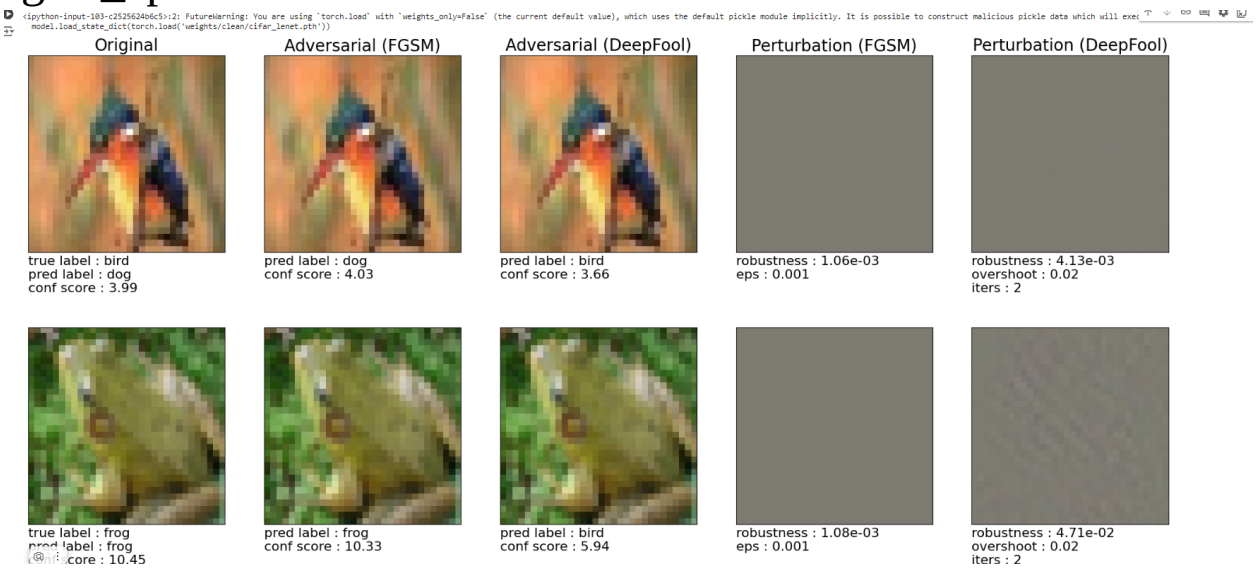

fgsm_eps = 0.02

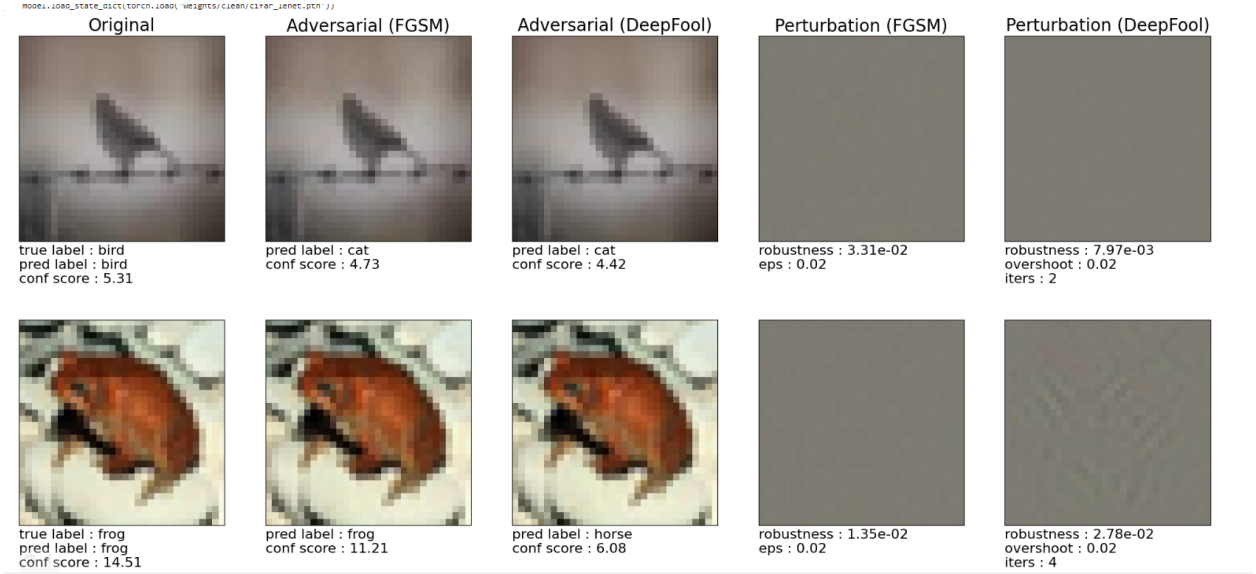

fgsm_eps = 0.5

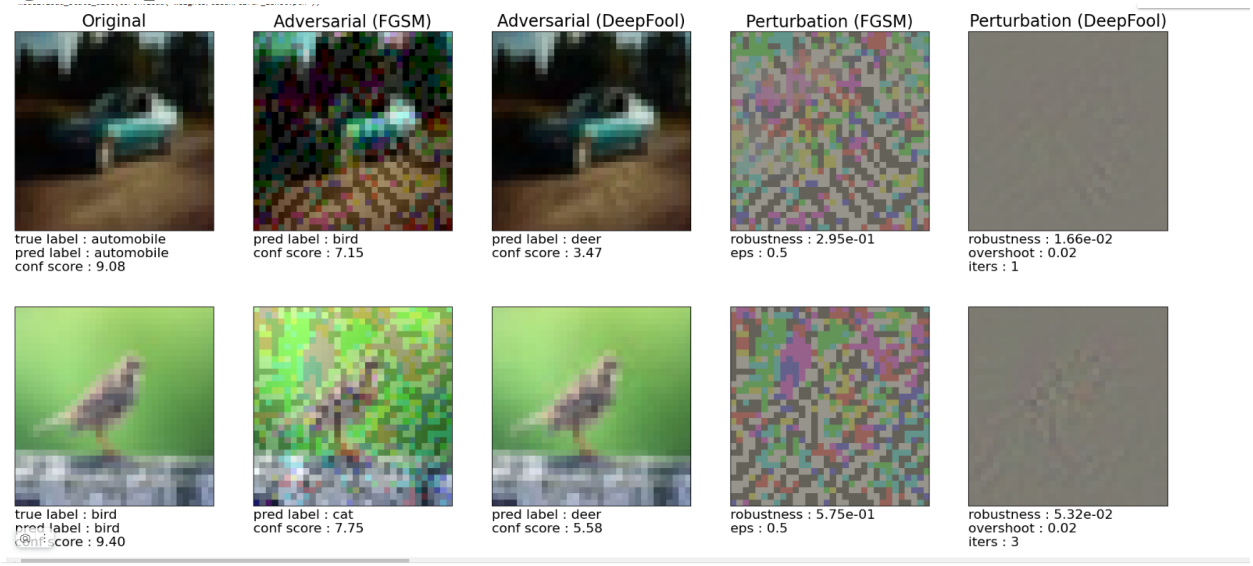

fgsm_eps = 0.9

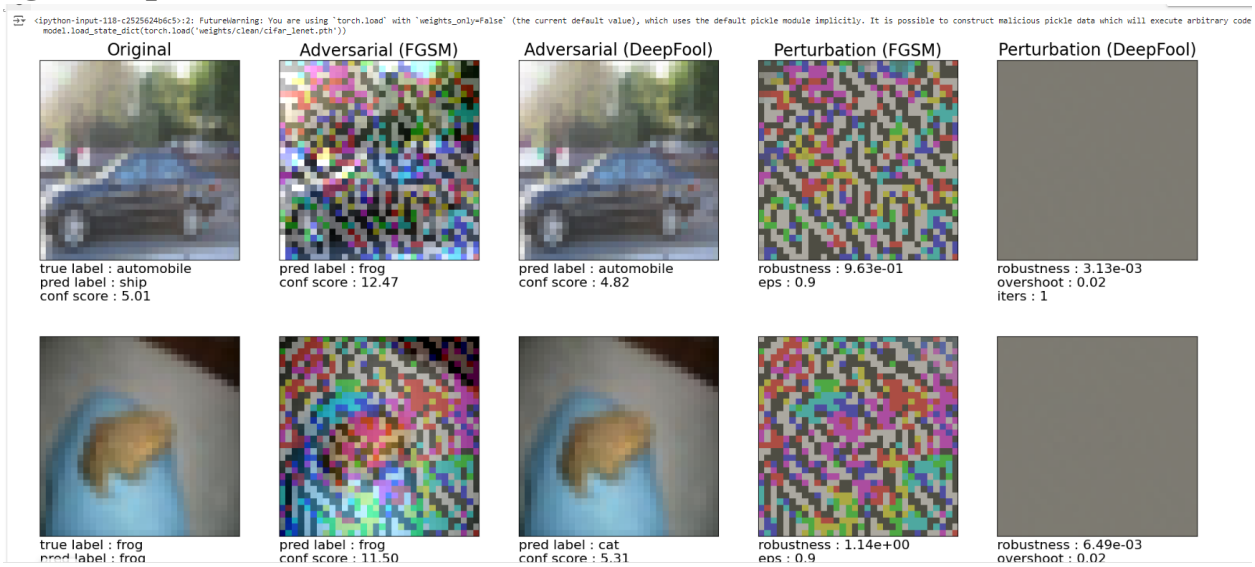

fgsm_eps = 10

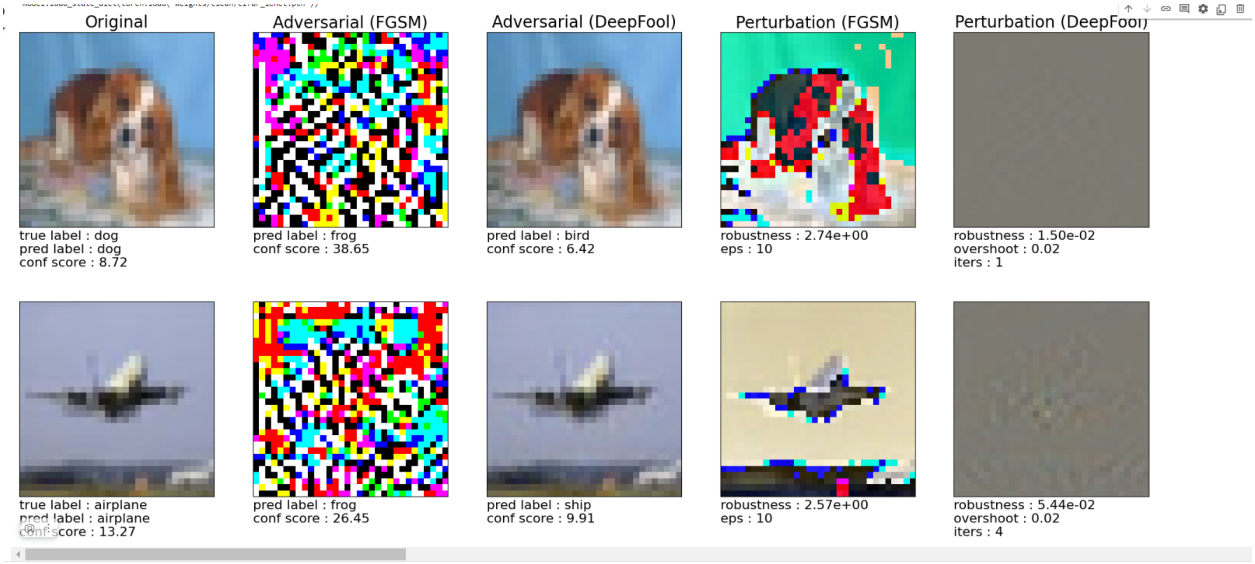

Заключение

Когда fgsm_eps увеличивается, сети становятся уязвимее к атакам. Значительно уязвимее они становятся со значения fgsm_eps = 0.5